

Steering undulatory microswimmers in a moving fluid through machine learning

Raphaël Chesneaux,¹ Laëtitia Giraldi,² and Jérémie Bec^{1,2}

¹*Ecole Nationale Supérieure des Mines de Paris, PSL University, CNRS, Cemef, Sophia-Antipolis, France*

²*Université Côte d’Azur, Inria, CNRS, Cemef, Sophia-Antipolis, France*

(Dated: July 27, 2021)

A new model of microswimmer is introduced to investigate locomotion strategies based on the sinusoidal undulation of a slender body. These active particles are then embedded in a prescribed fluid flow in which their swimming undulations have to compete with the drifts, strains and deformations inflicted by the outer velocity field. Such an intricate situation, where swimming and navigation are tightly bonded, is addressed using reinforcement learning. It is shown that the displacement strategy of the microswimmers can be efficiently optimized using Q -learning. Still, such an approach rapidly becomes rather expensive because of the highly chaotic character of the swimmers dynamics yielding a strong variability in learning efficiencies. A genetic algorithm is employed to circumvent such a pitfall.

I. INTRODUCTION

Microorganisms such as bacteria or plankton are natural examples of self-propelled particles. They often inspire the design of artificial devices used for industrial micro-manufacturing, toxic waste disposal, targeted drug delivery or localized medical diagnostics [?]. Recent technological developments on the use of micro-swimmers in medicine open new frontiers: those of surgery at the microscopic scale, directly inside the human body, enabling to deliver medicine and drugs in very precise places where their efficiency will be optimal. Much work was devoted to designing adequate nanorobots and studying the way they can be propelled and controlled using an external magnetic field [?], in particular for *in-vivo* conditions. Many questions are still open on how these micro-swimmers optimize their displacement, and in particular how they behave in complex flows comprising obstacles, walls, or having non-Newtonian properties. This is particularly important in order to find new designs and strategies that will allow artificial swimmers to reach new regions of the human body.

Studying and optimizing the displacement of swimmers and micro-swimmers is generally addressed in two successive steps. The first consists in finding an appropriate *swimming strategy* by choosing the composition, shape, or deformation that will lead to an efficient locomotion of the swimmer. The second step is to define a *navigation strategy* that will allow accounting for obstacles, for fluctuations in the surrounding flow, for its geometry, with the aim to minimize the time needed or the energy used to reach a target. Studying swimming strategies at the microscopic level requires advanced tools to describe fluid-structure interactions [? ?], to take a non-Newtonian rheology of the surrounding fluid into account [?], to model the hydrodynamics stresses due to the vicinity of walls [?]. Finding an effective strategy then relies on biomimetics [? ?] or on solving costly problems of optimal control [?]. As a matter of fact, such swimming issues are most of the time addressed in situations where the surrounding flow is at rest. This is

justified by the complexity and the computational costs that would be required to accurately model the intricate fluid-structure interactions occurring in a chaotic or turbulent medium.

Regarding navigation problems, there is an increasing interest in considering complicated carrier flows. The swimming mechanisms are then, most of the time, oversimplified and one rather focuses on how to adjust macroscopic active features of the swimmers in order to optimize their long-term displacement. Under such conditions, the use of machine learning techniques has proved efficiency [?]. Reinforcement learning has for instance been used to address navigation in a turbulent flow and to construct strategies that allow swimmers to find optimal paths to their targets in such a chaotic and fluctuating environment [? ? ?]. Navigation problems have also been studied from different perspectives such as finding new paths in the presence of obstacles that can be modelled as a potential barrier [?].

Here we want to address the two problems at once. The goal of this study is to demonstrate the applicability of machine learning approaches in a mesoscopic model of swimmer, and in particular to understand if such approaches are able, not only to make the swimmer move, but also to have it at the same time navigating in a complex environment. The swimmer will consist in a simple, deformable, inextensible thin filament described by the slender-body theory, and evolving in a space-varying incompressible fluid. Among the different types of swimming, we have chosen wave locomotion which is a self-propulsion strategy that relies on the generation and propagation of waves along the swimmer. This is a relatively simple, but remarkably robust technique that builds on the interactions between the swimmer and the fluid and appear in a variety of swimming strategies observed in nature. We will study the swimmer’s locomotion and navigation in a periodic cellular flow and then try to optimize its path using a Q -learning reinforcement algorithm.

II. A MATHEMATICAL MODEL OF MICROSWIMMER

A. Dynamics of a slender flagella

We consider that the swimmers are elongated, flexible and inextensible. They are moreover very thin, meaning that their cross-section diameter d is much smaller than their length ℓ . This assumption allows describing their interactions with the fluid in terms of the slender-body theory [?]. The swimmer is moreover embedded in an incompressible fluid flow whose unperturbed velocity field is denoted by $\mathbf{u}(\mathbf{x}, t)$. The swimmer's conformation at time t is characterized by a curve $\mathbf{X}(s, t)$ parametrized by its arc-length $s \in [-\ell/2, \ell/2]$. The dynamics of the fiber is given by the Cosserat equation, so that

$$\sigma \partial_t^2 \mathbf{X} = -\zeta \mathbb{R} [\partial_t \mathbf{X} - \mathbf{u}(\mathbf{X}, t)] + \partial_s (T \partial_s \mathbf{X}) - K \partial_s^4 \mathbf{X} + \mathbf{f}(s, t). \quad (1)$$

The first force on the right-hand side is the force exerted by the fluid. It involves the drag coefficient $\zeta = 8\pi\nu\rho_f/[2\log(\ell/d) - 1]$ (with ν the fluid kinematic viscosity and ρ_f its density) and the local Oseen's resistance tensor $\mathbb{R} = \mathbb{1} - (1/2) \partial_s \mathbf{X} \partial_s \mathbf{X}^\top$. The second force in the right-hand side is the tension whose amplitude T is determined by the inextensibility constraint $|\partial_s \mathbf{X}(s, t)| = 1$, valid at all time along the swimmer. The third term is the bending elasticity force which depends upon the swimmer's flexural rigidity K (product of Young's modulus and inertia). The last term denoted \mathbf{f} is an internal force that is prescribed to account for the so-called active behaviours of the swimmer and that is responsible for its locomotion. Equation (1) is associated with the free-end boundary conditions $\partial_s^2 \mathbf{X}(s, t) = 0$ and $\partial_s^3 \mathbf{X}(s, t) = 0$ at the fiber's extremities $s = \pm\ell/2$. The tension itself satisfies a second-order differential equation obtained by imposing $\partial_t |\partial_s \mathbf{X}|^2 = 0$ with the boundary conditions $T(s, t) = 0$ at $s = \pm\ell/2$.

In the limit of very small inertia, that is when the fiber's stopping time σ/ζ is much shorter than any characteristic time associated to the outer fluid velocity field \mathbf{u} , the dynamics further simplifies. It follows the overdamped limit obtained by balancing the viscous drag to the other forces in the slender-body equation, so that

$$\zeta \mathbb{R} [\partial_t \mathbf{X} - \mathbf{u}(\mathbf{X}, t)] = \partial_s (T \partial_s \mathbf{X}) - K \partial_s^4 \mathbf{X} + \mathbf{f}(s, t). \quad (2)$$

We focus on this equation in the sequel. The problem depends on several non-dimensional parameters that can be there identified. A first parameter is given by the ratio ℓ/L_f between the fiber's length ℓ and the spatial scale characteristic L_f of the fluid flow. Another parameter is $K/(\zeta L_f^3 U)$ where U is the order of magnitude of the fluid velocity. It measures the fiber's flexibility. The smaller it is, the more deformable is the fiber when it is subject to a compression or a shear by the flow. These two parameters are fixed for the rest of the study with representative values. We next turn to the undulatory models used to prescribe the force \mathbf{f} .

B. Model for undulatory locomotion

1. Helicoidal swimming

In order to better understand how to choose the force of locomotion, let us start with considering the case where there is no flow ($\mathbf{u} = 0$). We want to find under which conditions there exists a solution of (2) corresponding to the propagation of a wave along the fiber. We first choose this wave so that it has the shape of a circular helix and is associated with a displacement with a speed V along the axis of the helix. By placing ourselves in the reference frame $(\hat{x}, \hat{y}, \hat{z})$ where the axis of the helix is along \hat{z} , we are looking for a solution which, far from the edges, would take the form

$$\mathbf{X}(s, t) = \begin{bmatrix} R \cos(\nu s - \omega t) \\ \varepsilon R \sin(\nu s - \omega t) \\ V t + \tau s \end{bmatrix} \text{ where } R^2 \nu^2 + \tau^2 = 1 \quad \text{and } \varepsilon = \pm 1, \quad (3)$$

where R is the radius of the helix, $\tau \in [-1, 1]$ its torsion and ε its chirality. The curvature is constant, equal to $R\nu^2$, and the last condition corresponds to imposing the inextensibility of the fiber. It should be noted that this solution is not compatible with the boundary conditions satisfied by the solutions of (2). For now, it doesn't matter as we are looking for solutions taking the above form far enough from the ends of the fiber. By injecting (3) into the dynamical equation (2), we get

$$f_{\hat{x}} = K \nu^4 R \cos(\nu s - \omega t) + \zeta \left[\frac{\tau \nu}{2} V + \left(1 - \frac{R^2 \nu^2}{2} \right) \omega \right] R \sin(\nu s - \omega t), \quad (4)$$

$$f_{\hat{y}} = -\zeta \left[\frac{\tau \nu}{2} V + \left(1 - \frac{R^2 \nu^2}{2} \right) \omega \right] \varepsilon R \cos(\nu s - \omega t) + K \nu^4 \varepsilon R \sin(\nu s - \omega t), \quad (5)$$

$$f_{\hat{z}} = \zeta \left(1 - \frac{\tau^2}{2} \right) V + \zeta \frac{\tau \nu}{2} R^2 \omega, \quad (6)$$

The inextensibility condition $|\partial_s \mathbf{X}|^2 = 1$ is satisfied by the solution (3) at every moment of time. This explains why the terms involving the tension do not appear in the above system.

In the case where $\zeta \neq 0$ and where we impose that the force is not a global source of momentum for the swimmer, namely $\int \mathbf{f} ds = 0$, we get from (4-6) that ν has to be chosen as a multiple of $(2\pi/\ell)$ and that $f_{\hat{z}} = 0$. With this choice, the displacement speed is written as a function of the parameters of the locomotion force as

$$V = -\frac{\omega R^2 \nu \tau}{2 - \tau^2} = -\frac{\omega}{\nu} \frac{\tau(1 - \tau^2)}{2 - \tau^2}. \quad (7)$$

The swimming speed V in the absence of fluid flow is thus proportional to the phase velocity ω/ν with a constant that solely depends upon the helix torsion τ imposed by the locomotion forcing.

For our study, we decided to characterize the force amplitude by a non-dimensional parameter α , so that we have for \mathbf{f}

$$f_{\hat{x}} = B \cos(\nu s - \omega t) + A \sin(\nu s - \omega t), \quad (8)$$

$$f_{\hat{y}} = -A \cos(\nu s - \omega t) + B \sin(\nu s - \omega t), \quad (9)$$

$$f_{\hat{z}} = 0, \quad (10)$$

Where $A = \alpha \zeta \omega / \nu$ and $B = K \nu^4 R$. This solution therefore makes it possible to construct an adequate force for the locomotion of the fiber in a given direction \mathbf{p} . It is parameterized by the wave number ν , the frequency ω and the dimensionless parameter α . Its expression is of the form $\mathbf{f} = [\mathbf{m}, \mathbf{n}, \mathbf{p}] [f_{\hat{x}}, f_{\hat{y}}, 0]^T$ where \mathbf{m} and \mathbf{n} are two orthogonal unit vectors, both orthogonal to \mathbf{p} , the expressions of $f_{\hat{x}}$ and $f_{\hat{y}}$ being given by (8) and (9).

2. Planar undulatory swimming

Another choice of force that we have studied consists in propagating a sinusoidal plane wave along the fiber. This is a classical swimming strategy used by a sperm cell when it is trapped near a planar boundary surface [?]. In that case, the force is chosen to be

$$\mathbf{f}(s, t) = A \cos(\nu s - \omega t) \mathbf{p} \quad (11)$$

where \mathbf{p} is a unit vector in a direction orthogonal to that in which the swimmer is expected to move. Like in the helicoidal case, we write the forcing amplitude as $A = \alpha \zeta \omega / \nu$ where α is a dimensionless parameter that controls its strength. Furthermore we have the same constraint on ν coming from the global conservation of momentum.




FIG. 1. Swimming speed V as a function of the forcing amplitude parameter α for different wavenumbers k with $\nu = 2\pi k/\ell$ and a fixed value of the phase velocity ω/ν .

Contrarily to the helicoidal case, there is now no straightforward analytic relation between the swimming speed V and the force amplitude. This is due to the intricate role played by inextensibility and tension that prevents from obtaining an explicit solution for the fiber's conformation \mathbf{X} . We thus have recourse to numerics in

order to study the dependence of V upon the forcing parameter α . We expect from dimensional considerations on space and time scales that the swimming speed will again be proportional to the phase velocity ω/ν . This is confirmed by numerics. The dependence upon the amplitude parameter α and the wavenumber k is shown for a fixed value of ω/ν in Fig. 1. The swimming speed has a maximum (of the order of $0.2\omega/\nu$) that slowly increases as a function of k and shifts toward larger values of α . Our results suggests that achieving a given swimming speed becomes more energetic when the wave number of the undulation increases.

These numerical results were obtained by integrating the over-damped Cosserat equation (2) for isolated fibers in a fluid flow at rest. We use the second-order finite-difference scheme of [?] with $N = 201$ gridpoints along the fiber's arclength. The inextensibility constraint is achieved by a penalization method. Time marching uses a second-order semi-implicit Adams-Bashforth method with time step $\delta t = 5 \times 10^{-4}$.

We hereafter focus on this planar undulatory locomotion model with a wavenumber $k = 2$ and get interested in the following problem: we want the swimmer to swim as fast as possible in the $x > 0$ direction. This amounts to choose \mathbf{p} perpendicular to \mathbf{e}_x . It is interesting to point out the symmetry $x \mapsto -x$ of the force because it shows that forcing the fiber to swim in the \mathbf{e}_x direction by binding an \mathbf{e}_y force does not lead to a movement toward $x > 0$. In fact, depending on the orientation of the fiber, it can swim toward $x > 0$ or $x < 0$. This will have an impact on the naive swimming technique that is described below.

III. SWIMMING AND NAVIGATING IN A CELLULAR FLOW

From now on, we focus on the case of planar undulatory motility and consider that the swimmer is embedded in a two-dimensional cellular flow. More specifically, we choose to write the fluid velocity as $\mathbf{u} = (-\partial_y \Psi, \partial_x \Psi)$ with the stream function taking the simple periodic form $\Psi(x, t) = (LU/\pi) \cos(\pi x/L) \cos(\pi y/L)$. The spatial domain is hence covered by a tile of cells mimicking eddies. Their size L is chosen of the same order of magnitude as the fiber length ℓ . The velocity field has an amplitude U to be compared to the swimming velocity V introduced in previous section. Such a two-dimensional flow is a stationary solution of the incompressible Euler equations and is used to model the convection cells present in steady Rayleigh-Bénard convection. It is often employed to study the effects of fluid shear and rotation on transport and mixing. It moreover has the convenience of being easily reproducible by experiments [?]. Even if the motion of tracers in such a flow is non-chaotic, the swimmer's dynamics is. Our aim is here to study how a swimming fiber embedded in such a cellular flow is able

to move depending on various strategies. We first propose a naive strategy whose efficiency is then compared to that obtained from machine learning.

A. A naive strategy

Our aim is to maximize the swimmer's displacement toward the $x > 0$ direction. When using the basic swimming, that is to say always binding the fiber to swim with the force (11) constantly applied along the direction $\mathbf{p} = \mathbf{e}_y$, one obtains a completely chaotic motion. The fiber changes direction randomly: As illustrated in Fig. 2, when arriving in a new cell, either the swimmer passes through the cell, or its orientation is reversed and it leaves in the opposite direction. Consequently, the swimmer performs a random walk on long timescales. The average displacement of the fiber's center of mass is therefore zero.




FIG. 2. Case of a swimmer that continuously swims in the horizontal direction, with no specific strategy. It is initially released from the origin (shown as a cross) and the trajectory of its center of mass is shown as a gray curve. Several configurations are shown at different times. The colored background shows the fluid streamlines with, in red, the clockwise vortices and, in blue, those turning in the opposite direction.

A strategy allowing the swimmer to move in the $x > 0$ direction is what we call the *naive strategy*. It is simple: If the swimmer has the proper orientation, that is when the abscissa difference between its head and its center of mass is positive, the sinusoidal force (11) is applied in the direction $\mathbf{p} = \mathbf{e}_y$. If the fiber is wrongly aligned and faces the $x < 0$ direction, then we don't impose any force and the locomotion is stopped. This is a simple strategy that can be easily understood and induces a positive drift by breaking the symmetry $x \mapsto -x$. The results obtained with such a swimming are better than in the absence of any strategy and we get a positive average displacement.

We simulated around 50 fibers with different initial conditions. Each fiber starts at the center of a cell, with a completely unfold configuration and with a random angle with the horizontal, chosen between $-\pi/2$ and $\pi/2$. As we can see in Fig. 3 the naive strategy has a positive aver-




FIG. 3. Horizontal displacement of swimmers following the naive strategy, together with the average shown as a bold solid line, and the interval defined by the standard error shown as dashed lines. Time is here expressed in units of $(2\pi\omega)^{-1}$ and displacement in units of the cell size L .

age displacement but we observe that the distribution of swimmers significantly spreads with time. The standard deviation of the displacement increases with time and, in some cases, the displacement stays at an approximately constant value during rather long times, showing that the swimmers can get trapped in a cell and stop progressing to positive x . To improve our results and allow for a faster displacement, we next rely on machine-learning techniques and implement a Q -learning method.

B. An optimal strategy using Q -learning

We use the classical Q -learning algorithm with a set of discrete states \mathcal{S} and discrete actions \mathcal{A} . The policy is described by the Q -table, which assigns at time t to each couple $(s, a) \in \mathcal{S} \times \mathcal{A}$ a mark $Q_t(s, a)$. The action a_t taken at time t is such that it maximizes $Q_t(s_t, \cdot)$.

We define the swimmer's state as a combination of 3 features. The first property is the horizontal component of the fluid velocity at the fiber's head $u_h = \mathbf{e}_x \cdot \mathbf{u}(\mathbf{X}(0, t), t)$. We use a parameter u_0 and distinguish between three cases, $u_h < -u_0$, $u_h > u_0$, and $|u_h| < u_0$, which correspond to having a headwind, a tailwind, or no significant wind caused by the fluid flow. The second state parameter is the swimmer orientation, either towards positive x or towards negative x . It is defined as the sign of the abscissa difference between the fiber's head and its center of mass. The last state parameter accounts for the presence or not of a loop along the fiber. This last condition comes from observations of the fiber's swimming simulations where the buckling of the swimmer's body sometimes leads to a dead end, the fiber been locked in a given cell. All these parameters define a discrete set of 12 states. Concerning the actions, we define two categories: first the direction in which the force (11) is applied, either toward $\mathbf{p} = \mathbf{e}_x$ or toward $\mathbf{p} = \mathbf{e}_y$. Second the value of the amplitude parameter α that is chosen

in $\{0, 0.5, 0.75, 1\}$. There is, at the end, a discrete set of 8 actions.

The learning algorithm is based on the classical value-iteration update of the Bellman equation, namely

$$Q_{t+\Delta t}(s_t, a_t) = (1 - \lambda \Delta t) Q_t(s_t, a_t) + \lambda \Delta t \left[r_t + \gamma \max_a Q_t(s_{t+\Delta t}, a) \right], \quad (12)$$

where λ denotes the learning rate, γ the discount factor, and r_t is the reward, which in our case is defined as the distance travelled in the x direction by the swimmer's center of mass, namely $r_t = \bar{x}(t + \Delta t) - \bar{x}(t)$ with $\bar{x}(t) = (1/\ell) \int e_x \cdot \mathbf{X}(s, t) ds$. The time discretization of (12) involves a time step that we choose to be $\Delta t = 200 \delta t$. The learning rate is fixed to $\lambda = 0.05$, approximately corresponding to the inverse of the time needed by the swimmer to travel across two cells. Similarly, the discount factor is chosen such that the discount time $-\Delta t/(\log \gamma)$ corresponds to the time needed by the swimmer to traverse approximately ten cells. This weights the rewards obtained by the swimmer in the distant future, allowing our reinforcement strategy to optimize long-term displacements.

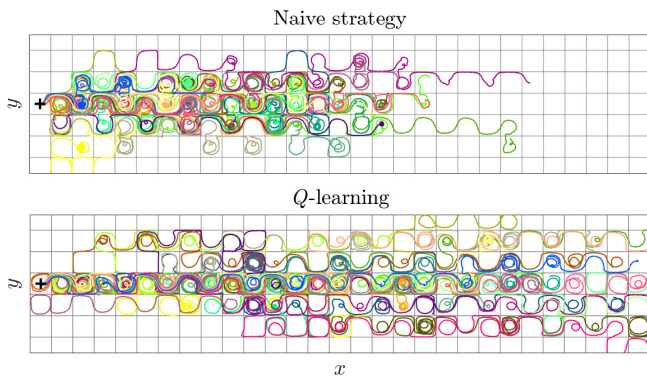


FIG. 4. Typical trajectories (various colors) obtained with the naive strategy (top) and with Q -learning (bottom). All swimmers are released from the origin (shown as a black cross) and try to move toward $x > 0$ across the various cells of our model flow.

We have simulated a set of 50 learning swimmers using for each of them different starting conditions randomly chosen. The entries of the Q -table are initialised to sufficiently large values, that is with optimistic premises, allowing the learning algorithm for wider explorations. Some trivially inefficient configurations have however been initially penalized. They correspond to situations where the fiber is wrongly oriented and swims against the fluid current, the later being strong enough to push the swimmer to the right. It is in that case much more efficient not to swim. We compare the performance of our Q -learning method with the naive strategy explained above. The Q -learning method rapidly display better results, as seen in Fig. 4. In the naive strategy, some fibers get locked in a cell and can not escape from it. This due

to the deformation of the fiber that can buckle the swimmer's body. We can also observe that the Q -learning method, in addition to locking less fiber in cells, allow them to go further and to progress further to the right.

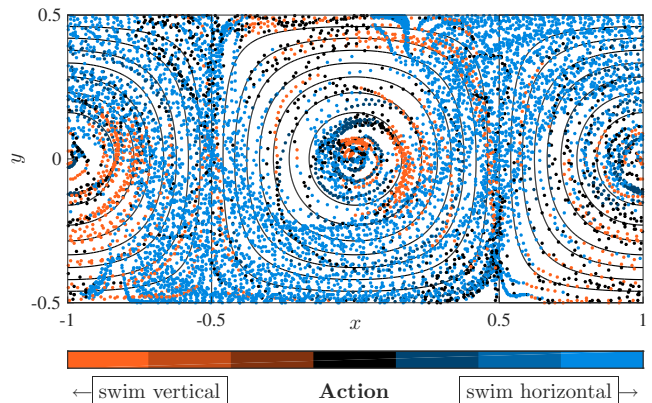


FIG. 5. Typical trajectory of the swimmer's center of mass obtained after learning. The positions have been here folded to account for periodicity. The colors correspond to the different actions taken by the swimmer, ranging from red (swimming with $\alpha = 1$ in the y direction) to blue (swimming with $\alpha = 1$ in the x direction).

Q -learning improves the naive strategy by giving a different path to the swimmers. With this method, the fiber takes advantage of the flow by going trough high positive horizontal fluid velocity area, passing alternately through the top of a cell and through the bottom. Such a “surfing” is evidenced from Fig. 5 that shows the typical trajectory of a swimmer after it has learnt an optimal strategy. The color code display the various actions that are performed by the swimmer. When surfing on an optimal path, the fiber maximizes its swim in the x direction. Conversely, when it get traps in the center of a cell or in a region of counterflow (*e.g.* top of the center cell), it either stops swimming or initiate a vertical displacement.




FIG. 6. Horizontal displacement for different realizations of learning, together with the average shown as a bold solid line, and the interval defined by the standard error shown as dashed lines.

Even if performance is significantly improved, we find that, as in the case of the naive swimmer, there is a rather strong variability in the displacement efficiency. This is evidenced from Fig. 6 that shows how a set of 50 trajectories spreads while it learns. Even if the average displacement outperforms by almost a factor 2 that obtained with the naive strategy in the same fluid flow (compare to Fig. 6), one observes that there is still a very strong variability among the various realizations. A large number of trajectories are very far from the average. Some, maybe lucky, perform much better. Others spend a very long time trapped in a cell (some of these events are visible in the bottom panel of Fig. 4). When trapped, they possibly forget a large part of the training they acquired before. Such a variability originates from the chaotic behavior of the fibers dynamics and is responsible for also a strong variability in the learning process.

C. Combining reinforcement learning with a genetic algorithm

To select efficient configurations and speed up the convergence of learning, we augment our Q -learning procedure using a genetic algorithm. For that, we put several fibers at competition during a time T and compare, at the end, the progression of each fiber to find the fastest one. The Q table of this fiber is then distributed to all the other fibers and then a learning run is started again for a time T with random initial conditions for each fiber. This approach is called a genetic algorithm because we only keep the learning part of the strongest fiber, as this is the case in nature.




FIG. 7. Implementation of the genetic algorithm applied here to five generations and a population of eight individuals. The horizontal displacement of individuals is periodically evaluated (at times multiple of $T = 500$, shown as vertical dotted lines) and only the strategy developed by the best candidate (blue curve) is retained and imposed to the next generation. Swimmers position are then reinitialised at the center of a cell with a random orientation. Such an algorithm accelerates the convergence to an asymptotic swimming speed V_{learn} (represented as a black dashed line).

This genetic algorithm requires more computing power but it significantly strengthens the convergence of learning towards an optimal swimming strategy. In fact, as the learning takes place, there is a convergence towards a single value V_{learn} of the average speed of movement of the fibers. In fact, as shown in Fig. 7, as the various runs are carried out, we note that all realizations get gradually concentrated close to each other, at variance with what was previously obtained. The dispersion of the results decreases during the different learning periods. This highlights the convergence of learning toward an optimal strategy that allows to obtain reproducible results.

The simulations reported in Fig. 7 were carried out for a fixed value of the fluid flow amplitude U . In order to demonstrate the performance of learning using genetic algorithm, we carried out simulations by varying the value of this parameter. This allows us to underline the dependence of the efficiency of learning on the dimensionless parameter U/V defined as the ratio between this amplitude and the swimming speed V in absence of any fluid flow. We can see in Fig. 8 that at all values of this parameter, Q -learning leads to a more efficient strategy than the naive one. This improvement is the most efficient when the ratio U/V is close to 1. In fact, in this case, the influence of the fluid makes it possible for swimming strategy to have an impact. In the case when $U/V \ll 1$, the fluid has no impact on the swim, a straight swim is most effective. Finally in the case $U/V \gg 1$, it is the fluid which is the most important and the choice of the strategy does not matter (the fiber undergoes the flow).




FIG. 8. Asymptotic velocity of the swimmers (here normalized by the swimming speed V_{swim} in the absence of fluid flow), as a function of the fluid flow amplitude U , with both the average value obtained with the naive strategy and the optimal value obtained with Q -learning, as labeled.

For the naive strategy, we can see that the travelling velocity has a minimum for $U/V \approx 1$. This is due to the fact that when the fluid amplitude is equal to the swimming speed, the fiber can be trapped in a region where it experiences a headwind that exactly matches its displacement. Such a trapping disappears when $U/V > 1$, because the fluid wins and is able to push the fiber in a

more favorable position. Finally, we have shown that the learning method that we implemented can work under different speed regimes although it can be more efficient in certain specific cases.

IV. SUMMARY AND PERSPECTIVES

To conclude, we have presented here a mathematical models to study simultaneously the swim and the navigation of a slender fiber in a two-dimensional cell flow. We have developed a Q -learning approach to optimize the displacement strategy of the fiber and maximize its velocity in a prescribed direction. We have also devel-

oped a genetic algorithm that, combined with Q -learning, can converge to an optimal strategy with better velocities than those from a naive strategy. We have therefore shown the relevance of the use of Q -learning in learning to swim in a flow displaying spatial variations.

A possible opening for future work would consist in studying the movements of a fiber, that have been previously trained, in more chaotic turbulent flows. It is indeed a possibility that would allow studying the effects of learning with the Q -learning method in a more realistic situation. In addition, the next step could also consist in improving the learning method by using reinforcement deep learning methods with neural networks. All this makes it possible to consider a vast field of study on the motility and navigation of such microswimmers.